



CS6650

Cache Money



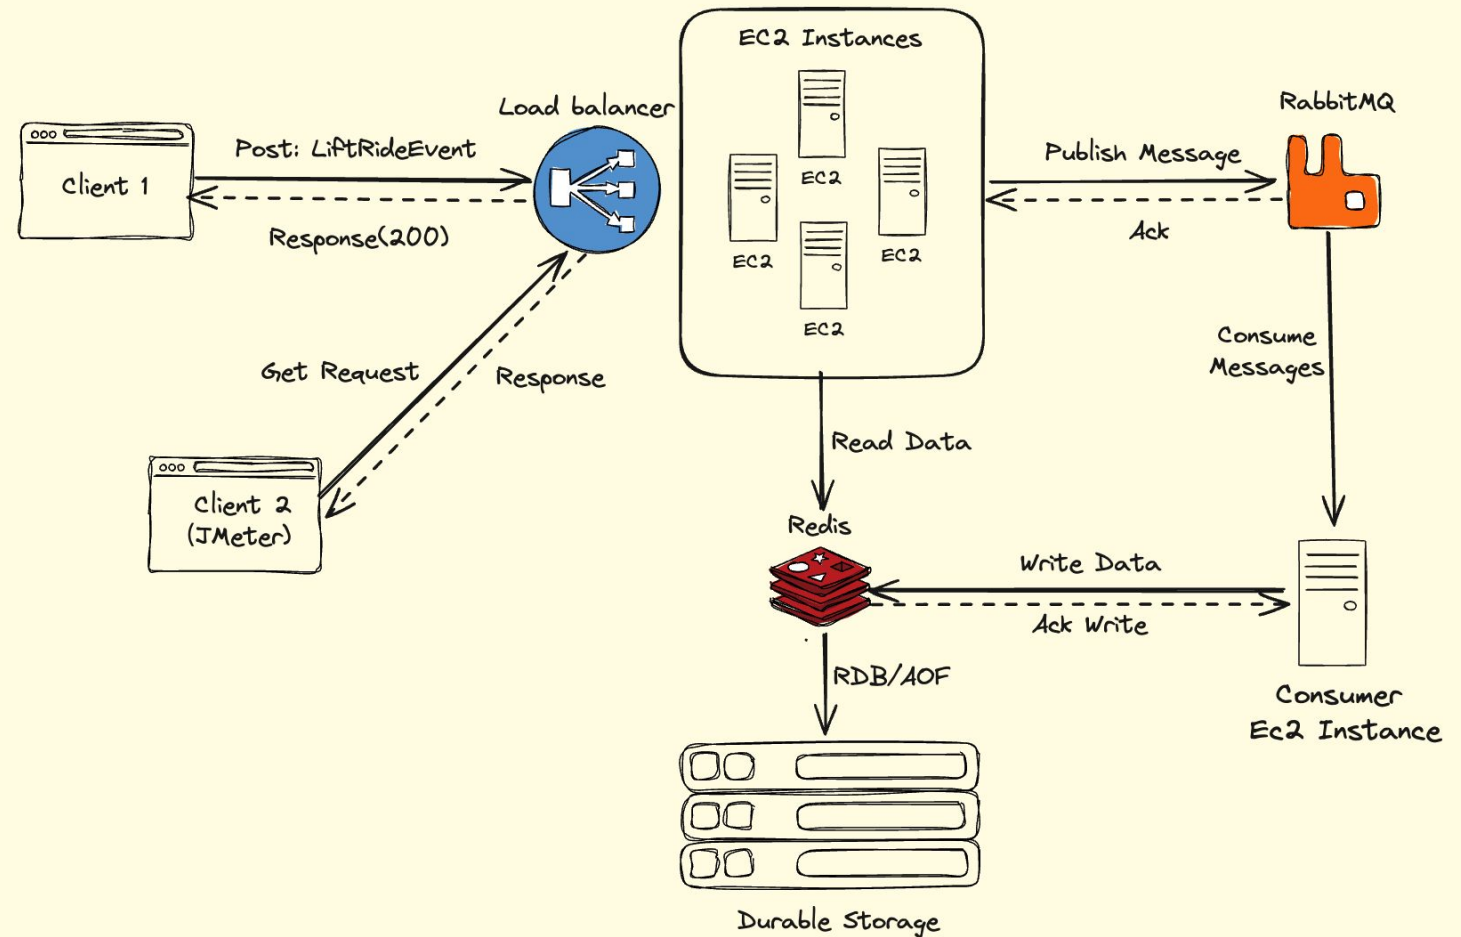
CONTENTS

- Architecture
- Data model
- Jmeter tests
- Future improvement

Architecture Overview

- **Application Load Balancer** with a total of 4 EC2 Instances
- **RabbitMQ** deployed on a single instance for message queuing and async processing
- **Redis** used as a database, with RDB and AOF for durability
- **Consumer** on a separate instance to consume messages from MQ and writes to Redis.
- **Client 1** sends 'LiftRideEvent' POST requests to server
- **Client 2** sends GET request for load testing (JMeter)

System Architecture



DATA MODEL

- We use three independent data structures in Redis, each corresponding to a separate GET request.

- GET/resorts/{resortID}/seasons/{seasonID}/day/{dayID}/skiers (Redis Set)

SET	resort:9:season:2024:day:1:skiers	No limit	818 KB
-----	-----------------------------------	----------	--------

- GET/skiers/{resortID}/seasons/{seasonID}/days/{dayID}/skiers/{skierID} (Redis String)

STRING	skier:66430:resort:10:season:2024:day:1	No limit	88 B
--------	---	----------	------

- GET/skiers/{skierID}/vertical (Redis String)

STRING	skier:45757:vertical	No limit	64 B
--------	----------------------	----------	------

TRADE OFF

- Pros of Using Redis:
 1. Extremely High Performance
 2. Simple Data Structures
 3. Scalability: Redis can be scaled easily by increasing the number of nodes in the cluster.
- Cons of Using Redis
 1. Data Durability
 2. Memory Limitation

DISASTER RECOVERY

- Enable RDB in disk snapshot in redis.conf
- Redis RDB persistence is a snapshotting method where Redis saves its dataset to disk at specified intervals. This approach is beneficial for disaster recovery because it provides a point-in-time backup of data, which can be used to restore the system after a crash.
- Using RDB will lose some performance, but considering the disaster recovery and availability improvements, this is a reasonable trade off.

```

[ec2-user@ip-172-31-29-157 ~]$ redis-cli KEYS '*' | wc -l
267782
[ec2-user@ip-172-31-29-157 ~]$ sudo kill -9 $(pgrep redis-server)
[ec2-user@ip-172-31-29-157 ~]$ sudo systemctl status redis
● redis.service - Redis In-Memory Data Store
   Loaded: loaded (/etc/systemd/system/redis.service; enabled; vendor preset: disabled)
   Active: failed (Result: signal) since Wed 2024-04-17 21:48:17 UTC; 4s ago
     Process: 2962 ExecStart=/usr/local/bin/redis-server /etc/redis/redis.conf (code=killed, signal=KILL)
    Main PID: 2962 (code=killed, signal=KILL)

Apr 17 21:44:47 ip-172-31-29-157.us-west-2.compute.internal systemd[1]: Started Redis In-Memory Data..
..
Apr 17 21:48:17 ip-172-31-29-157.us-west-2.compute.internal systemd[1]: redis.service: main process ..
.L
Apr 17 21:48:17 ip-172-31-29-157.us-west-2.compute.internal systemd[1]: Unit redis.service entered f..
[..
Apr 17 21:48:17 ip-172-31-29-157.us-west-2.compute.internal systemd[1]: redis.service failed.
Hint: Some lines were ellipsized, use -l to show in full.
[ec2-user@ip-172-31-29-157 ~]$ sudo reboot
Connection to ec2-52-12-252-4.us-west-2.compute.amazonaws.com closed by remote host.
Connection to ec2-52-12-252-4.us-west-2.compute.amazonaws.com closed.
zhangyunfei@zhangyunfeiMacBook-Pro ~ % ssh -i "/Users/zhangyunfei/Downloads/key2.pem" ec2-user@ec2-52-
12-252-4.us-west-2.compute.amazonaws.com
Last login: Wed Apr 17 21:46:35 2024 from 23.252.51.83

      #_
     ~\  #####_      Amazon Linux 2
    ~ ~ \_#####\
    ~ ~  \###|      AL2 End of Life is 2025-06-30.
    ~ ~   \#/  ---
    ~ ~    V~'  '--->
    ~ ~      /
    ~ ~     /      A newer version of Amazon Linux is available!
    ~ ~    _/_/_/
    ~ ~   _/_/_/      Amazon Linux 2023, GA and supported until 2028-03-15.
    ~ ~  _/m/'         https://aws.amazon.com/linux/amazon-linux-2023/

[
8 package(s) needed for security, out of 17 available
Run "sudo yum update" to apply all updates.
[ec2-user@ip-172-31-29-157 ~]$ redis-cli KEYS '*' | wc -l
267782
[ec2-user@ip-172-31-29-157 ~]$ █

```

JMETER TESTS

- The JMeter tests were conducted on an EC2 instance, employing two different server configurations: a single server and a load-balanced setup with 4 instances. Each server configuration was evaluated using three different thread counts: 128, 300, and 400.
- Test with 128 threads gave us the lowest response times and p99s.

JMETER tests results: GET/resorts/{resortID}/seasons/{seasonID}/day/{dayID}/skiers

Single service

Statistics

Requests	Executions			Response Times (ms)							Throughput	Network (KB/sec)	
Label ^	#Samples ^	FAIL ^	Error % ^	Average ^	Min ^	Max ^	Median ^	90th pct ^	95th pct ^	99th pct ^	Transactions/s ^	Received ^	Sent ^
Total	64000	0	0.00%	8.87	1	79	8.00	16.00	19.00	24.00	4680.42	795.56	735.89
Unique-Day	64000	0	0.00%	8.87	1	79	8.00	16.00	19.00	24.00	4680.42	795.56	735.89

Load balance

Statistics

Requests	Executions			Response Times (ms)							Throughput	Network (KB/sec)	
Label ^	#Samples ^	FAIL ^	Error % ^	Average ^	Min ^	Max ^	Median ^	90th pct ^	95th pct ^	99th pct ^	Transactions/s ^	Received ^	Sent ^
Total	64000	0	0.00%	7.22	1	74	8.00	12.00	16.00	24.00	5033.03	738.84	938.78
Unique-Day	64000	0	0.00%	7.22	1	74	8.00	12.00	16.00	24.00	5033.03	738.84	938.78

JMETER tests results: GET/skiers/{resortID}/seasons/{seasonID}/days/{dayID}/skiers/{skierID}

Single service

Statistics

Requests	Executions			Response Times (ms)							Throughput	Network (KB/sec)	
Label ^	#Samples ⇅	FAIL ⇅	Error % ⇅	Average ⇅	Min ⇅	Max ⇅	Median ⇅	90th pct ⇅	95th pct ⇅	99th pct ⇅	Transactions/s ⇅	Received ⇅	Sent ⇅
Total	64000	0	0.00%	17.50	2	1011	16.00	17.00	24.00	25.00	3495.17	491.51	947.05
Get Request 2	64000	0	0.00%	17.50	2	1011	16.00	17.00	24.00	25.00	3495.17	491.51	947.05

Load balance

Statistics

Requests	Executions			Response Times (ms)							Throughput	Network (KB/sec)	
Label ^	#Samples ⇅	FAIL ⇅	Error % ⇅	Average ⇅	Min ⇅	Max ⇅	Median ⇅	90th pct ⇅	95th pct ⇅	99th pct ⇅	Transactions/s ⇅	Received ⇅	Sent ⇅
Total	64000	0	0.00%	8.81	1	58	8.00	16.00	18.00	24.00	4734.08	798.46	1144.06
Get Request 2	64000	0	0.00%	8.81	1	58	8.00	16.00	18.00	24.00	4734.08	798.46	1144.06

JMETER tests results: GET/skiers/{skierID}/vertical

Single service

Statistics

Requests	Executions			Response Times (ms)							Throughput	Network (KB/sec)	
Label	#Samples	FAIL	Error %	Average	Min	Max	Median	90th pct	95th pct	99th pct	Transactions/s	Received	Sent
Total	64000	0	0.00%	8.78	1	307	8.00	16.00	18.00	24.00	4681.10	797.38	653.20
Skier-verticals	64000	0	0.00%	8.78	1	307	8.00	16.00	18.00	24.00	4681.10	797.38	653.20

Load balance


Statistics

Requests	Executions			Response Times (ms)							Throughput	Network (KB/sec)	
Label	#Samples	FAIL	Error %	Average	Min	Max	Median	90th pct	95th pct	99th pct	Transactions/s	Received	Sent
Total	64000	0	0.00%	7.18	1	69	8.00	13.00	16.00	20.00	4963.93	730.66	838.10
Skier-verticals	64000	0	0.00%	7.18	1	69	8.00	13.00	16.00	20.00	4963.93	730.66	838.10

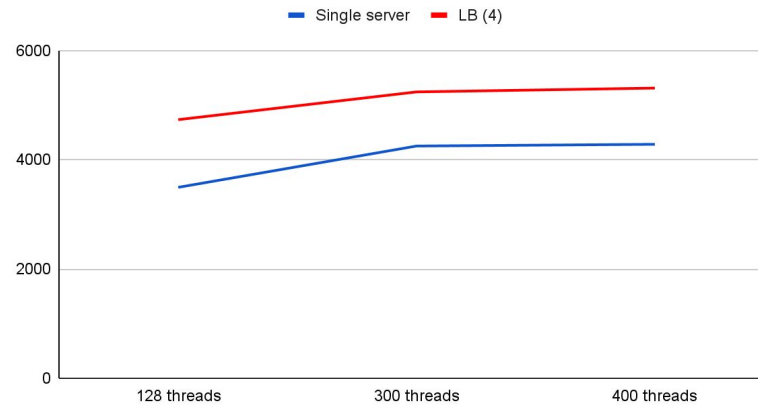
JMETER tests results

Single service VS Load balance


Throughputs(/s)



Throughputs(/s)




Throughput(/s)



Response time(ms)

AVG	128 threads	300 threads	400 threads
Single server	8.87	37.81	57.12
LB (4)	7.22	22.44	36.59

Future improvement



- Current Design (store all the data in the redis)
 - Occupied much memory
 - Can't support other type of query very easily
- Further Design (use MySQL to store data, use ES to speed up query)
 - Use MySQL to store origin data (skier's data)
 - Use Redis as cache to speed up the GET query
 - Use Elasticsearch to Speed up complex search (join search) and aggregation